

UC Berkeley  
Teaching Professor  
Dan Garcia

# CS61C

## Great Ideas in Computer Architecture (a.k.a. Machine Structures)



UC Berkeley  
Professor  
Bora Nikolić

## Introduction to the C Programming Language

# Bugs, and Pointers

# C Syntax: Variable Declarations

- Similar to Java, but with a few minor but important differences
  - All variable declarations must appear before they are used
  - All must be at the beginning of a block.
  - A variable may be initialized in its declaration; *if not, it holds garbage!*
    - the contents are undefined...
- Examples of declarations:
  - Correct: { int a = 0, b = 10; ...
  - Incorrect in ANSI C: for (int i=0; ...
  - Correct in C99 (and beyond): for (int i=0; ...

# An Important Note: Undefined Behavior...

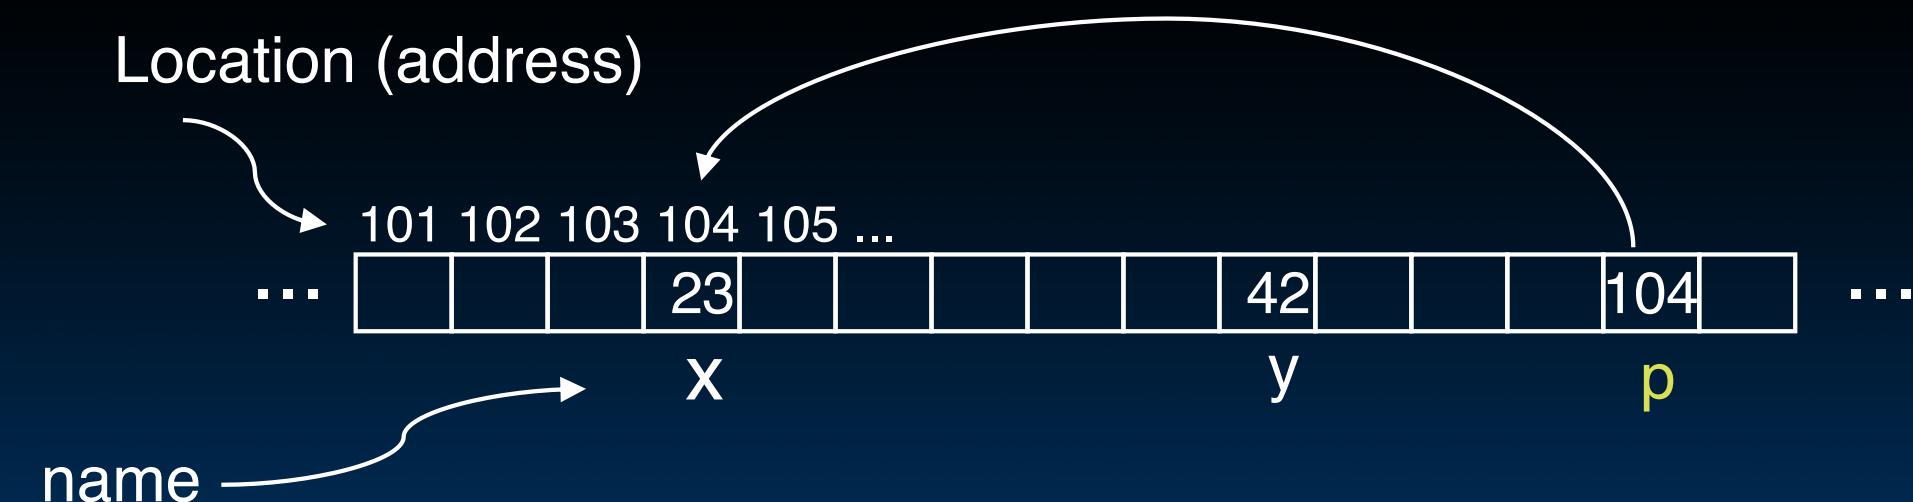
- A lot of C has “Undefined Behavior”
  - This means it is often unpredictable behavior
    - It will run one way on one computer...
    - But some other way on another
    - Or even just be different each time the program is executed!
- Often characterized as “Heisenbugs”
  - Bugs that seem random/hard to reproduce, and seem to disappear or change when debugging
  - Cf. “Bohrbugs” which are repeatable

# Address vs. Value

- Consider memory to be a single huge array:
  - Each cell of the array has an address associated with it.
  - Each cell also stores some value.
  - Do you think they use signed or unsigned numbers?  
Negative address?!
- Don't confuse the **address** referring to a memory location with the **value** stored in that location.
- For now, the abstraction lets us think we have access to  $\infty$  memory, numbered from 0...



- An address refers to a particular memory location. In other words, it points to a memory location.
- **Pointer:** A variable that contains the address of a variable.



# Pointer Syntax

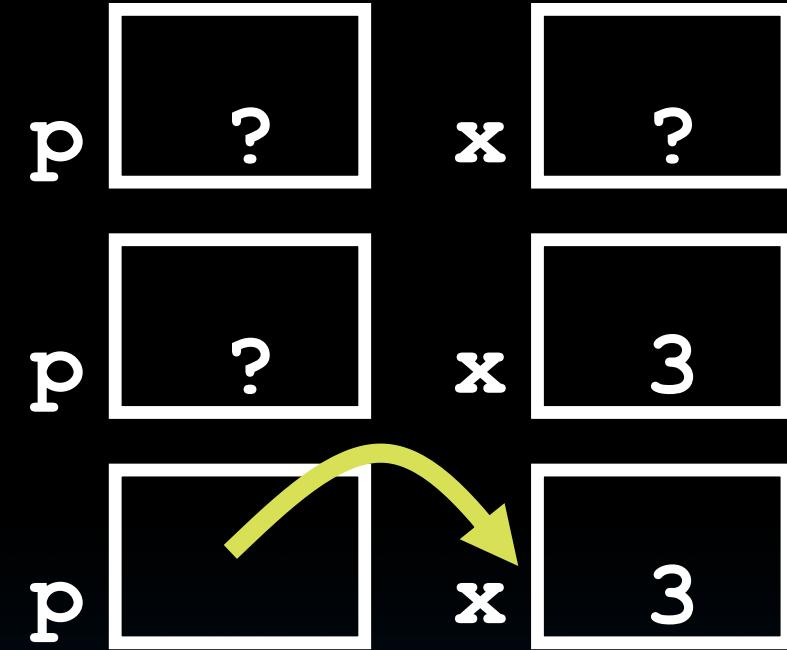
- `int *p;`
  - Tells compiler that variable **p** is address of an **int**
- `p = &y;`
  - Tells compiler to assign **address of y** to **p**
  - **&** called the “**address operator**” in this context
- `z = *p;`
  - Tells compiler to assign **value at address in p** to **z**
  - **\*** called the “**dereference operator**” in this context

- How to create a pointer:  
& operator: get address of a variable

```
int *p, x;
```

```
x = 3;
```

```
p = &x;
```

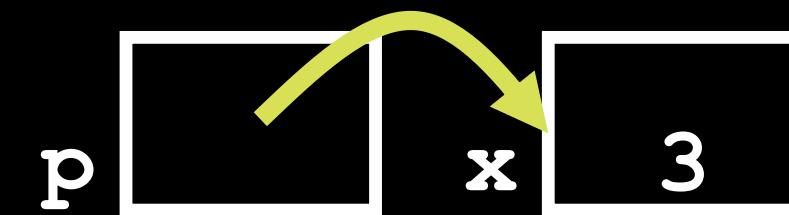


Note the “\*” gets used 2 different ways in this example. In the declaration to indicate that **p** is going to be a pointer, and in the **printf** to get the value pointed to by **p**.

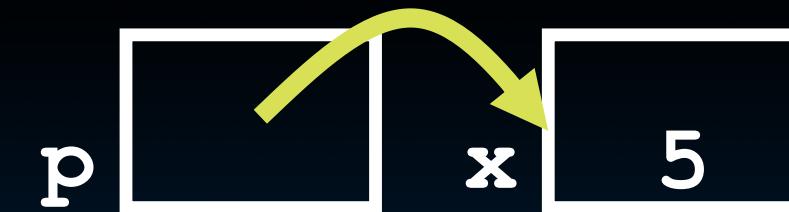
- How get a value pointed to?
  - \* “dereference operator”: get value pointed to

```
printf("p points to %d\n", *p);
```

- How to change a variable pointed to?
  - Use dereference \* operator on left of =



`*p = 5;`



# Pointers and Parameter Passing (1/2)

- **Java and C pass parameters “by value”**
  - procedure/function/method gets a copy of the parameter, so changing the copy cannot change the original

```
void addOne (int x) {  
    x = x + 1;  
  
}  
  
int y = 3;  
addOne(y);
```

y is still = 3

# Pointers and Parameter Passing (2/2)

- How to get a function to change a value?

```
void addOne (int *p) {  
    *p = *p + 1;  
}  
  
int y = 3;  
addOne (&y);
```

y is now = 4

# More C Pointer Dangers

- Declaring a pointer just allocates space to hold the pointer – it does not allocate something to be pointed to!
- Local variables in C are not initialized, they may contain anything.
- What does the following code do?

```
void f()
{
    int *ptr;
    *ptr = 5;
}
```

# Pointers in C... The Good, Bad, and the Ugly

- **Why use pointers?**
  - If we want to pass a large struct or array, it's easier / faster / etc. to pass a pointer than the whole thing
    - Otherwise we'd need to copy a huge amount of data
  - In general, pointers allow cleaner, more compact code
- **So what are the drawbacks?**
  - Pointers are probably the single largest source of bugs in C, so be careful anytime you deal with them
    - Most problematic with dynamic memory management—coming up next time
    - Dangling references and memory leaks



# Using Pointers Effectively

- Pointers are used to point to **any** data type (**int**, **char**, a **struct**, etc.).
- Normally a pointer can only point to one type (**int**, **char**, a **struct**, etc.).
  - **void \*** is a type that can point to anything (generic pointer)
  - Use sparingly to help avoid program bugs... and security issues... and a lot of other bad things!
- You can even have pointers to functions...
  - `int (*fn) (void *, void *) = &foo`
    - **fn** is a function that accepts two **void \*** pointers and returns an **int** and is initially pointing to the function **foo**.
    - `(*fn)(x, y)` will then call the function

# Pointers and Structures

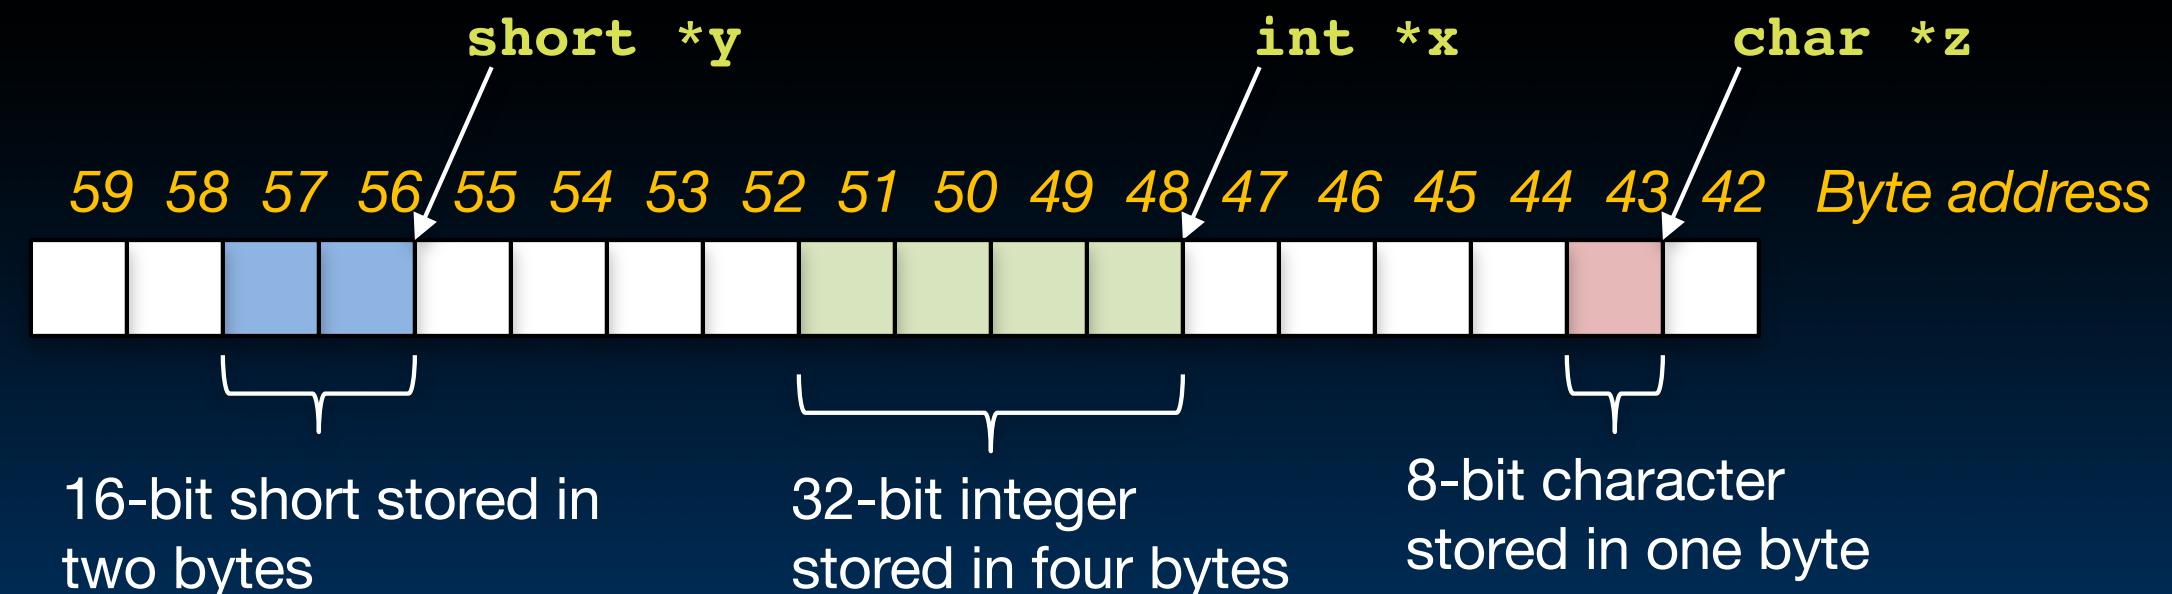
```
typedef struct {  
    int x;  
    int y;  
} Point;  
  
Point p1;  
Point p2;  
Point *paddr;  
  
/* dot notation */  
int h = p1.x;  
p2.y = p1.y;  
  
/* arrow notation */  
int h = paddr->x;  
int h = (*paddr).x;  
  
/* This works too */  
p1 = p2;
```

# NULL pointers...

- **The pointer of all 0s is special**
  - The "NULL" pointer, like in Java, python, etc...
- **If you write to or read a null pointer, your program should crash**
- **Since "0 is false", its very easy to do tests for null:**
  - `if(!p) { /* P is a null pointer */ }`
  - `if(q) { /* Q is not a null pointer */ }`

# Pointing to Different Size Objects

- Modern machines are “byte-addressable”
  - Hardware’s memory composed of 8-bit storage cells, each has a unique address
- A C pointer is just abstracted memory address
- Type declaration tells compiler how many bytes to fetch on each access through pointer
  - Eg., 32-bit integer stored in 4 consecutive 8-bit bytes
- But we actually want “word alignment”
  - Some processors will not allow you to address 32b values without being on 4 byte boundaries
  - Others will just be very slow if you try to access “unaligned” memory.



# Arrays

# Arrays (1/5)

- **Declaration:**
  - `int ar[2];`
  - ...declares a 2-element integer array
  - An array is really just a block of memory
- **Declaration and initialization**
  - `int ar[] = {795, 635};`
  - declares and fills a 2-elt integer array
- **Accessing elements:**
  - `ar[num]`
  - returns the num<sup>th</sup> element.

# Arrays (2/5)

- **Arrays are (almost) identical to pointers**
  - `char *string` and `char string[]` are nearly identical declarations
  - They differ in very subtle ways: incrementing, declaration of filled arrays
- **Key Concept: An array variable is a “pointer” to the first element.**

# Arrays (3/5)

- **Consequences:**
  - `ar` is an array variable but looks like a pointer in many respects (though not all)
  - `ar[0]` is the same as `*ar`
  - `ar[2]` is the same as `* (ar+2)`
  - We can use pointer arithmetic to access arrays more conveniently.
- **Declared arrays are only allocated while the scope is valid**

```
char *foo() {  
    char string[32]; ...;  
    return string;  
} is incorrect
```

# Arrays (4/5)

- Array size **n**; want to access from 0 to **n-1**, so you should use counter AND utilize a variable for declaration & incr
  - Wrong

```
int i, ar[10];
for(i = 0; i < 10; i++){ ... }
```
  - Right

```
int ARRAY_SIZE = 10;
int i, a[ARRAY_SIZE];
for(i = 0; i < ARRAY_SIZE; i++){ ... }
```
- Why? **SINGLE SOURCE OF TRUTH**
  - You're utilizing **indirection** and **avoiding maintaining two copies** of the number 10

# Arrays (5/5)

- **Pitfall: An array in C does not know its own length, & bounds not checked!**
  - Consequence: We can accidentally access off the end of an array.
  - Consequence: We must pass the array and its size to a procedure which is going to traverse it.
- **Segmentation faults and bus errors:**
  - These are **VERY** difficult to find; be careful!
  - You'll learn how to debug these in lab...

# Pointer Arithmetic

- **pointer + n**
  - Adds **n\*sizeof** (“whatever pointer is pointing to”) to the memory address
  
- **pointer – n**
  - Adds **n\*sizeof** (“whatever pointer is pointing to”) to the memory address

# Pointers (1/4) ...review...

- **Java and C pass parameters “by value”**
  - procedure/function/method gets a copy of the parameter, so changing the copy cannot change the original

```
void addOne (int x) {  
    x = x + 1;  
  
}  
  
int y = 3;  
addOne(y);
```

y is still = 3

# Pointers (2/4) ...review...

- How to get a function to change a value?

```
void addOne (int *p) {  
    *p = *p + 1;  
}  
  
int y = 3;  
  
addOne (&y);
```

y is now = 4

# Pointers (3/4)

- But what if you want to change a pointer?
  - What gets printed?

```
void IncrementPtr(int *p)
{    p = p + 1; } *q = 50

int A[3] = {50, 60, 70};
int *q = A;
IncrementPtr( q );
printf("*q = %d\n", *q);
```

The diagram illustrates the state of memory after the execution of the `IncrementPtr` function. It shows an array `A` with three elements: 50, 60, and 70. A pointer `q` initially points to the first element (50). After the call to `IncrementPtr(q)`, the value of `q` is modified to point to the second element (60). The variable `*q` is also updated to reflect this new value.

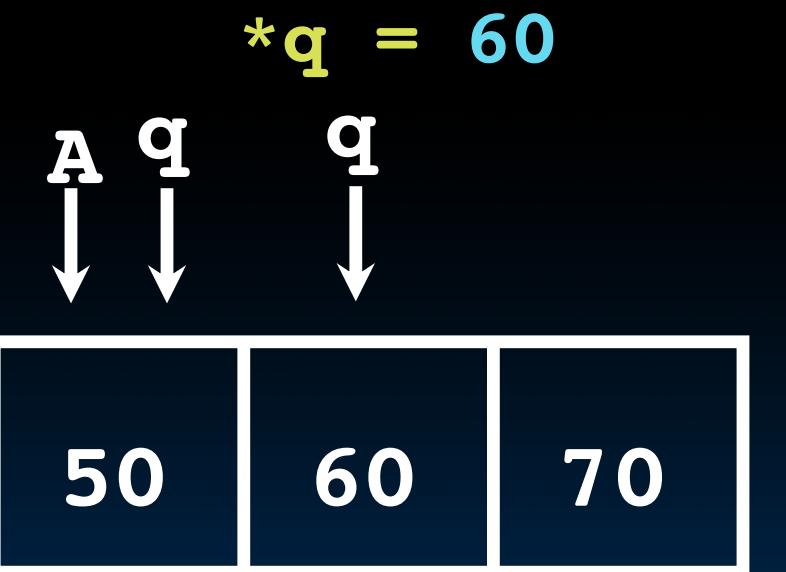
# Pointers (4/4)

- Idea! Pass a pointer to a pointer!
  - Declared as  $\star\star h$
  - Now what gets printed?

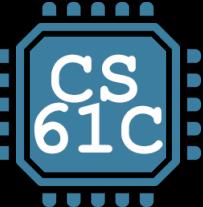
```
void IncrementPtr(int **h)
{
    *h = *h + 1; }
```

```
int A[3] = {50, 60, 70};
int *q = A;
IncrementPtr(&q);
printf("*q = %d\n", *q);
```



# Function Pointer Example



# map (actually `mutate_map` easier)

```
#include <stdio.h>

int x10(int), x2(int);
void mutate_map(int [], int n, int(*)(int));
void print_array(int [], int n);

int x2 (int n) { return 2*n; }
int x10(int n) { return 10*n; }

void mutate_map(int A[], int n, int(*fp)(int)) {
    for (int i = 0; i < n; i++)
        A[i] = (*fp)(A[i]);
}

void print_array(int A[], int n) {
    for (int i = 0; i < n; i++)
        printf("%d ",A[i]);
    printf("\n");
}
```

```
% ./map
3 1 4
6 2 8
60 20 80
```

```
int main(void)
{
    int A[] = {3,1,4}, n = 3;
    print_array(A, n);
    mutate_map (A, n, &x2);
    print_array(A, n);
    mutate_map (A, n, &x10);
    print_array(A, n);
}
```